Ruth Contract Playbook

Version 1.1

Technology Services & AI/ML Development Engagements

Fixed-Price and Time & Materials Contracts with Global Clients

Prepared for:	Ruth – Commercial, Legal & Delivery Teams
Company Profile:	150-person technology services company Full-stack development and ML/AI development
Document Status:	Internal – Confidential
Supersedes:	Ruth Contract Playbook – Version 1.0
Scope:	Sections 2–15 cover Preferred Position, Acceptable Alternatives, High-Risk/Red Lines, and Negotiation Guidance for every major contract topic.

This playbook is an internal negotiation and risk-management tool. It does not constitute legal advice and does not override written approval requirements set out in Ruth's Delegation of Authority policy.

Contents

1 PURPOSE, SCOPE AND HOW TO USE THIS PLAYBOOK

Purpose. This playbook guides Ruth's commercial, legal, and delivery teams in negotiating technology services and AI/ML development contracts consistently. It is designed to protect Ruth's intellectual property, manage liability exposure appropriate to a 150-person services firm, and enable fast, confident deal closure with global clients across fixed-price and time-and-materials (T&M) engagements.

Scope. This playbook applies to all Master Service Agreements (MSAs), Statements of Work (SOWs), Order Forms, Data Processing Addenda (DPAs), and client-authored paper submitted for full-stack development and ML/AI development engagements, regardless of deal size, client geography, or commercial model.

Structure of Sections 2–15. Each substantive topic follows a consistent four-part structure:

- **Preferred Position** – Ruth's ideal opening ask, stated in specific, deployable contract language.
- **Acceptable Alternatives** – fallback positions that deal teams may agree to *without further escalation*, subject to the stated conditions.
- **High-Risk / Red Lines** – positions that require Legal and/or Leadership sign-off before acceptance. Each item is labeled **[High Risk]** or **[Medium Risk]** . **[High Risk]** items default to *rejection*; any exception requires documented senior approval before signature. **[Medium Risk]** items may proceed only with a documented mitigation (e.g., pricing adjustment, insurance increase, narrower scope).
- **Negotiation Guidance** – practical talking points and trade-offs a commercial lead can use directly in a live negotiation.

Escalation path. Any **[High Risk]** item requires written approval from the Head of Legal (or General Counsel, once appointed) and, where liability exposure, IP transfer, or committed pricing exceeds the thresholds defined in Ruth's Delegation of Authority policy, sign-off from the CEO or COO, before any signature is executed.

What's new in Version 1.1. Version 1.1 expands Version 1.0 with materially deeper AI/ML-specific provisions across Intellectual Property (Section 5), Warranties (Section 7), Limitation of Liability (Section 8), Confidentiality & Data Protection (Section 6), Service Levels (Section 9), and Acceptance Testing (Section 10). These additions reflect Ruth's growing AI/ML development practice and its increasing exposure to foundation-model dependencies, training-data provenance issues, model performance expectations, and AI-specific regulatory scrutiny from global clients. Version 1.1 also adds explicit treatment of open-source and AI tooling obligations (Section 13), key-person and staffing clauses (Section 12), audit rights (Section 14), and most-favored-nation pricing and force majeure (Section 15).

Who should use this. Sales/commercial leads drafting or reviewing SOWs, delivery leads flagging operational risk in client paper, and Legal reviewing MSAs. Any deviation beyond an Acceptable Alternative must be routed through Legal using the escalation path above – this playbook does not authorize deal teams to concede High-Risk items unilaterally.

2 PARTIES, CONTRACT STRUCTURE AND PRECEDENCE OF DOCUMENTS

Preferred Position

- Ruth contracts through the appropriate Ruth legal entity for the client's region; the relationship is structured as an MSA plus one or more SOWs, where the SOW governs scope, price, schedule, and deliverables, and the MSA governs legal terms (IP, liability, confidentiality, termination).
- Order of precedence for conflicts: (1) SOW-specific commercial terms (scope, price, schedule) control over the MSA for those specific items; (2) the MSA's IP, liability, confidentiality, and indemnification terms always control over any conflicting SOW or client purchase-order language, without exception.
- Ruth engages as an independent contractor. The agreement expressly disclaims any partnership, joint venture, agency, or employment relationship between the parties.

Acceptable Alternatives

- A single combined MSA-plus-SOW document (no separate MSA) is acceptable for engagements under USD 50,000 in total contract value or under 3 months' duration.
- Client-provided paper may serve as the base agreement provided Ruth attaches a signed addendum that expressly overrides the client paper's IP, liability, payment, and confidentiality sections with Ruth's standard terms.
- Reversing precedence so that client paper controls over Ruth's MSA is acceptable only for enterprise clients with over USD 500,000 in annual contract value, and only after full Legal review and sign-off.
- Multi-entity contracting (e.g., a client subsidiary signs while the parent company provides a payment/performance guarantee) is acceptable provided the guarantee is captured in writing as part of the signed agreement.

High-Risk / Red Lines

- **[High Risk]** Client paper's terms govern the entire relationship with no Ruth addendum – this exposes Ruth to unknown or unlimited liability, undefined IP treatment, and unfavorable payment terms, and must not be signed without Legal review.
- **[High Risk]** Engagement proceeds on verbal or email-only scope with no signed SOW – no delivery work should begin without a written, signed SOW.
- **[Medium Risk]** Evergreen auto-renewal of the MSA with no defined notice period, or a notice period shorter than 30 days for non-renewal.
- **[Medium Risk]** Contract language names individual Ruth engineers (rather than the Ruth entity) as contracting parties, creating personal liability exposure for staff.
- **[High Risk]** Client contract requires Ruth's AI infrastructure or tooling subcontractors (e.g., cloud/AI model providers) to be directly and individually bound by the client's own contractual terms – this is not achievable, since Ruth has no privity to modify third-party providers' standard terms.

Negotiation Guidance

- Always open with Ruth's MSA as the base paper; if the client insists on their own paper, escalate to Legal before any commercial terms are finalized, not after.
- Negotiate scope and commercial terms in the SOW first, then reconcile legal terms in the

MSA in parallel – this avoids stalling delivery timelines over legal redlines.

- Offer the client a single "conformed" clean copy showing all agreed changes once redlines are resolved, to speed up final internal sign-off on the client side.
- If a client's procurement team insists on their own paper for policy reasons, propose a short Ruth-specific addendum (2–3 pages) covering only IP, liability cap, and payment terms, rather than trying to replace the entire document.

3 SCOPE OF WORK, STATEMENTS OF WORK AND CHANGE CONTROL

Preferred Position

- Every SOW defines: detailed scope and deliverables, acceptance criteria, a fixed-price milestone schedule (or T&M rate card with estimated hours), documented assumptions and dependencies, and an explicit out-of-scope list.
- All scope changes are handled through a written Change Order (CO), signed by both parties, before any additional work begins; the CO specifies the revised price, timeline, and any impact on deliverables.
- For AI/ML SOWs, target model performance metrics (accuracy, F1, latency, etc.) are stated as *development targets to be pursued using best professional efforts*, not as guaranteed outcomes, and the SOW explicitly names the source, format, and ownership of training data to be used.

Acceptable Alternatives

- Minor changes (under 8 hours of effort or under USD 2,000 in value) may be approved via email confirmation and formalized retroactively in the next invoice, rather than a full CO.
- For Agile/T&M engagements, a rolling backlog model is acceptable, where scope is confirmed sprint-by-sprint via a sprint planning document instead of a formal CO, provided the rate card and estimated team velocity are fixed in the SOW.
- A "not-to-exceed" band on fixed-price SOWs (e.g., up to 15% cost overage) may be absorbed without a new CO, billed at standard T&M rates, provided it is pre-agreed and capped.
- A client-authored scope document may be accepted as the base scope description provided Ruth attaches an assumptions-and-exclusions addendum before signature.

High-Risk / Red Lines

- **[High Risk]** Open-ended "best efforts" scope with no defined deliverables or acceptance criteria combined with a fixed-price commercial structure – this exposes Ruth to unbounded rework at a fixed fee.
- **[High Risk]** SOW commits to a specific AI model accuracy or performance threshold (e.g., "the model shall achieve 95% accuracy") as a binding contractual deliverable, with no carve-out for training data quality, availability, or real-world data drift.
- **[Medium Risk]** Client is granted a unilateral right to expand scope "within the existing budget," effectively enabling scope creep without a corresponding change to price or timeline.
- **[Medium Risk]** No change-control mechanism defined at all, such that any client email or verbal instruction is treated as a binding amendment to the SOW.
- **[High Risk]** SOW requires Ruth to build a model using client-supplied proprietary or third-party training data with no representations from the client as to the data's quality, lineage, consent basis, or legality.
- **[Medium Risk]** Fixed-price commercial structure applied to genuinely exploratory or R&D-stage ML work, where the outcome and required effort are inherently uncertain at the time of quoting.

Negotiation Guidance

- Push exploratory or model-training/R&D phases toward T&M or a T&M-with-cap structure; reserve fixed-price commercial terms for well-defined build and integration phases only.

- Consistently frame AI/ML performance commitments as "target metrics, pursued on a best-efforts basis, subject to agreed data quality" rather than guaranteed outcomes – this is a recurring point of client pushback and should be a standard talking point.
- Offer a defined change-order turnaround SLA (e.g., CO reviewed and returned within 3 business days) to make the change-control process feel low-friction to the client, reducing resistance to formal CO requirements.
- If a client refuses any change control mechanism, propose a monthly scope-reconciliation checkpoint as a lighter-weight compromise rather than accepting no process at all.

4 PRICING, PAYMENT TERMS AND INVOICING

Preferred Position

- Payment terms of Net 30 from invoice date; invoicing occurs monthly in arrears for T&M engagements and upon milestone completion for fixed-price engagements.
- Late payment accrues interest at 1.5% per month (or the highest rate permitted by applicable law, if lower), and Ruth may suspend work after payment is 15 days past due, with 5 business days' prior written notice of suspension.
- Rate card is fixed for the initial 12 months of the engagement; thereafter, rates may increase by up to 8% annually with 30 days' written notice to the client.

Acceptable Alternatives

- Net 45 payment terms for enterprise clients with over USD 1,000,000 in annual contract value and an established, on-time payment history with Ruth.
- A milestone holdback of up to 10% of fixed-price milestone value, released upon final acceptance, may substitute for an upfront deposit, provided the holdback is released within 90 days of delivery regardless of unrelated disputes.
- Where a client requires invoicing in local currency rather than USD, an FX adjustment clause tied to a defined reference rate (e.g., the rate on invoice date, with a stated tolerance band) may be added.
- Quarterly (rather than annual) rate review may be agreed for engagements expected to exceed 18 months in duration, provided any increase remains capped at the agreed annual ceiling.

High-Risk / Red Lines

- **[High Risk]** Payment to Ruth is made contingent on the client first receiving payment from its own end-customer ("pay-when-paid") – this shifts the client's counterparty credit risk onto Ruth.
- **[High Risk]** No contractual right for Ruth to suspend work for non-payment – this removes Ruth's principal commercial leverage in a payment dispute.
- **[Medium Risk]** Payment terms of Net 60 or longer with no corresponding late-payment interest or fee mechanism.
- **[Medium Risk]** Client may withhold payment of an *entire* invoice over a dispute regarding a single line item, rather than withholding only the disputed portion.
- **[High Risk]** Most-favored-nation (MFN) pricing clause requiring Ruth to automatically extend its lowest global rate for any client, anywhere, to this client – operationally unworkable across a multi-client services business and commercially damaging if agreed broadly.
- **[Medium Risk]** Outcome-based or success-fee pricing tied solely to AI model performance in production, where the client controls the production data pipeline, infrastructure, and deployment context that materially affect that performance.

Negotiation Guidance

- Trade longer payment terms (Net 45/60) for a larger upfront deposit (e.g., 20–30% of fixed-price value) to protect cash flow without conceding on headline payment days.
- If a client insists on an MFN clause, narrow it tightly to "same scope, same volume, same region, same commercial model" and cap the review frequency to once per year on written request only.
- Position the suspension-of-work right as a de-escalation tool in payment disputes (pause work,

resolve, resume) rather than an immediate termination trigger – this is usually more palatable to client procurement teams.

- If outcome-based pricing is commercially important to win the deal, condition it explicitly on client-provided data quality and infrastructure commitments documented in the SOW, so performance shortfalls attributable to the client do not reduce Ruth's fee.

5 INTELLECTUAL PROPERTY RIGHTS AND DATA/MODEL OWNERSHIP

Preferred Position

- Client owns all client-specific deliverables (custom application code, client-specific fine-tuned model weights, documentation, and designs created under the SOW), with IP assignment taking effect *upon receipt of full and final payment* for the applicable deliverable – not upon creation.
- Ruth retains ownership of all pre-existing and independently developed IP ("Ruth Background IP"), including general-purpose frameworks, internal ML libraries, accelerators, prompt-engineering templates, and reusable tooling, and grants the client a perpetual, worldwide, royalty-free license to use such Background IP solely as embedded in the delivered deliverables.
- For fine-tuned or custom-trained models built exclusively on the client's proprietary data for the client's exclusive use, ownership of the resulting fine-tuned model weights transfers to the client upon full payment; the underlying base/foundation model, model architecture, and training pipeline/tooling used to produce those weights remain the property of Ruth or its third-party licensor and are licensed, not assigned.

Acceptable Alternatives

- Joint ownership of genuinely novel, jointly-developed algorithms or inventions (with each party free to use the joint IP independently, royalty-free) may be agreed for strategic co-innovation engagements, but not as a default position.
- Ruth may retain a non-exclusive right to reuse general learnings, non-confidential techniques, and improved internal tooling developed during the engagement – excluding client data, client-specific model weights, and client confidential information – for use with other clients.
- A client reference/case-study right (client name plus a high-level, non-confidential description of the engagement) may be requested in exchange for concessions elsewhere, unless the client is designated as a competitor-sensitive account.
- For productized components intended for reuse across multiple clients, Ruth retains ownership of the component and grants the client a broad, non-exclusive commercial-use license instead of a full IP assignment.

High-Risk / Red Lines

- **[High Risk]** Client contract language claims ownership of Ruth's pre-existing frameworks, internal ML accelerators, or proprietary training pipelines used to deliver the engagement – this must always be excluded via a clearly defined Background IP carve-out.
- **[High Risk]** IP assignment takes effect "upon creation" rather than "upon full payment" – this removes Ruth's leverage if the client later disputes or withholds payment mid-project.
- **[High Risk]** Client requires a perpetual, royalty-free right to use Ruth's general-purpose base models or proprietary training techniques for unrelated future products, at no additional fee – this effectively transfers Ruth's core technology assets for the price of one engagement.
- **[Medium Risk]** Broad "work made for hire" language drafted broadly enough to potentially sweep in Ruth's internal tools and utilities that were not built specifically for this client.
- **[High Risk]** Client contract requires Ruth to indemnify for IP infringement in the resulting trained model while remaining silent on the provenance of the training data the client itself supplied – Ruth cannot control or verify data it did not source.
- **[Medium Risk]** No carve-out for open-source components embedded within the deliverables, creating a direct conflict with a "full and unencumbered IP assignment" warranty (see also

Section 13).

Negotiation Guidance

- Always maintain a clear, defined-terms distinction in the contract between "Deliverable IP" (assigned to the client) and "Background IP / Tools" (licensed, not assigned) – this single distinction resolves the majority of IP disputes before they arise.
- On model ownership specifically, propose the standard split: fine-tuned/custom weights trained on client data = client-owned; base model, architecture, and training methodology = Ruth or third-party-owned and licensed for the client's use.
- If a client insists on full ownership of Ruth's underlying tooling or platform, price that explicitly as a separate technology buyout fee, distinct from and additional to the services fee.
- Tie the effective date of any IP assignment to "full and final payment received" as a standing negotiation point – this is a legitimate, widely-accepted market practice and protects Ruth's position in payment disputes.

6 CONFIDENTIALITY AND DATA PROTECTION

Preferred Position

- Mutual, reciprocal confidentiality obligations with a 3-year survival period post-termination for general Confidential Information, and indefinite survival for trade secrets.
- Client Data and Confidential Information are used only for the purposes of the engagement; Ruth will not use client Confidential Information or Client Data to train Ruth's own general-purpose models, or share it with third-party AI model providers for model-training purposes, without the client's prior written consent.
- A Data Processing Addendum (DPA) is executed for any engagement involving personal data, specifying the parties' roles (controller/processor), a defined sub-processor list (including any AI infrastructure vendors), and an appropriate cross-border transfer mechanism (e.g., Standard Contractual Clauses) where applicable.

Acceptable Alternatives

- Confidentiality survival may be reduced to 2 years for non-sensitive technical information, provided trade secrets and personal data protections remain unaffected.
- A pre-approved sub-processor list with a 15-day objection window for new AI vendors may substitute for prior written consent on each individual addition, to keep delivery moving.
- Use of anonymized or aggregated client data for Ruth's internal benchmarking or tool improvement may be permitted with the client's explicit written opt-in, separate from the base confidentiality terms.
- Where a client requires zero third-party AI model provider exposure, Ruth may offer an on-premises or client-VPC-hosted open-weight model deployment as an alternative to a contractual data-use restriction, priced as an additional infrastructure line item.

High-Risk / Red Lines

- **[High Risk]** Blanket contractual prohibition on Ruth using any third-party AI tools (e.g., LLM-based coding assistants) anywhere in delivery, with no workable carve-out – this is operationally unworkable at scale and should be replaced with an approved-tools list mechanism.
- **[High Risk]** Client requires Ruth to *guarantee* zero data retention by third-party LLM/AI infrastructure providers, when Ruth has no technical or contractual ability to enforce that guarantee against the underlying provider.
- **[Medium Risk]** Unlimited confidentiality survival period with no standard exclusions for information that is independently developed, becomes public through no fault of Ruth, or was already known prior to disclosure.
- **[High Risk]** Client Data including personal data (PII) is transferred to Ruth with no DPA, no defined cross-border transfer mechanism, and no breach notification timeline specified.
- **[Medium Risk]** Client demands direct audit or inspection rights over Ruth's contracts with its AI vendors/sub-processors – rarely obtainable in practice and should be replaced with a summary attestation or compliance report.
- **[High Risk]** Breach notification period shorter than 48 hours combined with uncapped liability for the resulting breach – an operationally unrealistic and commercially dangerous combination.

Negotiation Guidance

- Maintain a standing "Approved AI Tools & Sub-processors" list as a schedule to the DPA, updatable with notice rather than case-by-case consent, to preserve delivery efficiency while giving the client visibility and control.

- Offer a standard 72-hour breach notification commitment (aligned with common GDPR practice) as Ruth's default position, rather than negotiating from an undefined starting point.
- If a client insists on zero third-party model exposure, present the private/VPC-hosted or self-hosted open-source model option as a priced alternative rather than declining outright.
- Where a client requests direct system or vendor-contract audit rights, counter-offer a SOC 2 report, security questionnaire response, or third-party attestation as the standard verification mechanism.

7 WARRANTIES

Preferred Position

- Ruth warrants that deliverables will materially conform to the SOW specifications for 90 days following acceptance (the "Warranty Period"), and that Ruth will re-perform or fix, at no additional charge, any material defect in the deliverables caused by Ruth's own work.
- For AI/ML deliverables, Ruth warrants only that the model was built using the agreed methodology and evaluated against the agreed test dataset as of the delivery date; Ruth does *not* warrant ongoing production accuracy, correctness of model outputs, or the absence of hallucination or bias, since real-world performance depends on post-deployment data drift and usage patterns outside Ruth's control.
- Ruth warrants that it has the right to grant the licenses granted under the agreement, and that, to its knowledge, the deliverables do not knowingly infringe any third party's intellectual property rights – excluding client-supplied data, third-party APIs, foundation models, and open-source components used as specified or approved by the client.

Acceptable Alternatives

- The Warranty Period may be extended to 180 days for fixed-price engagements exceeding USD 250,000, in exchange for a defined paid maintenance retainer commencing immediately thereafter.
- Warranty remedy may be limited to re-performance only (no refund), unless re-performance fails to cure the defect after two attempts, at which point a pro-rata refund applies only to the fees for the specific defective component.
- A performance-based warranty (e.g., "the model will meet the agreed benchmark score on the agreed test set as of delivery") may be agreed only where the client supplies, and both parties sign off on, the evaluation dataset and methodology within the SOW itself.
- Warranty exclusions may be extended to explicitly cover client-side modifications, third-party integrations not built by Ruth, or use of the deliverables outside documented parameters.

High-Risk / Red Lines

- **[High Risk]** Warranty of a specific ongoing AI model accuracy or error rate *in live production* (e.g., "the model shall maintain 98% accuracy at all times") – Ruth cannot control post-deployment data drift, client data pipeline changes, or downstream misuse.
- **[High Risk]** Uncapped warranty duration (e.g., "warranted for the life of the software" or with no stated end date).
- **[Medium Risk]** Warranty extends to third-party components (open-source libraries, third-party APIs, foundation models) beyond a simple pass-through of the original vendor's own warranty, if any.
- **[High Risk]** Implied warranties of merchantability and fitness for a particular purpose are not expressly disclaimed – this creates open-ended exposure under the default commercial law of many jurisdictions.
- **[Medium Risk]** Warranty remedy includes consequential or indirect losses arising from the defect, rather than being limited to repair, re-performance, or a capped refund.
- **[High Risk]** Client requires a warranty against "bias," "hallucination," or "discriminatory outcomes" of an AI model with no defined, agreed test methodology or dataset – this is an unmeasurable and effectively unbounded obligation.

Negotiation Guidance

- Clearly separate the "delivery-time conformance warranty" (bounded, testable, time-limited) from "ongoing production performance" (explicitly disclaimed), and offer an optional paid MLOps/monitoring retainer to cover the latter as a value-add rather than a free obligation.
- Where a client raises bias or fairness concerns, propose a defined, mutually-agreed evaluation protocol and dataset upfront, rather than accepting an open-ended warranty against subjective outcomes.
- Always pair any warranty grant with an explicit disclaimer of all implied warranties permitted by applicable law (merchantability, fitness for purpose, non-infringement beyond the stated scope).
- Position an extended warranty period as a paid add-on or bundled with a support retainer, rather than a free concession, to protect margin during negotiation.

8 LIMITATION OF LIABILITY AND INDEMNIFICATION

Preferred Position

- Aggregate liability of either party under the agreement is capped at the total fees paid by the client in the 12 months preceding the claim under the applicable SOW (or the total SOW value, whichever is greater, for engagements under 12 months).
- Mutual exclusion of indirect, incidental, consequential, special, and punitive damages, and loss of profits, revenue, or business, for both parties – except for breaches of confidentiality, IP infringement, or gross negligence/willful misconduct, which are carved out from this exclusion but remain subject to the liability cap.
- Ruth indemnifies the client only for third-party claims that Ruth's original deliverables (excluding client-supplied data, open-source components, and third-party/foundation-model components) infringe a third party's intellectual property rights, subject to the standard liability cap; the client indemnifies Ruth for claims arising from client-supplied data, specifications, or instructions.

Acceptable Alternatives

- The liability cap may be raised to 1.5x–2x total fees for enterprise engagements exceeding USD 500,000, in exchange for higher insurance requirements and a corresponding pricing premium.
- A separate, higher carve-out cap (e.g., 3x fees or a defined fixed dollar amount) may apply specifically to confidentiality or data-breach claims, distinct from the general liability cap.
- A "cap-within-cap" structure may be used for IP indemnification specifically, if the client insists on separate treatment from general liability – but this must remain a defined, finite figure, never uncapped.
- For data-breach-heavy engagements, the liability cap may instead be tied to the limits of Ruth's applicable cyber/E&O insurance policy, rather than to contract fees.

High-Risk / Red Lines

- **[High Risk]** Uncapped liability of any kind, including for AI model output errors, data breaches, or IP infringement – this is a fundamental red line regardless of deal size.
- **[High Risk]** Ruth is required to indemnify the client for any harm caused by AI model outputs in production (e.g., discriminatory decisions, incorrect medical or financial recommendations) regardless of root cause – this improperly shifts business and regulatory risk onto Ruth for a system whose real-world behavior depends heavily on the client's own data and deployment context.
- **[Medium Risk]** The liability cap explicitly excludes "any breach of the agreement" as a carve-out category, which in practice uncaps liability for ordinary contractual breaches.
- **[High Risk]** Client requires Ruth to indemnify for third-party IP claims arising from the underlying foundation model or LLM provider (e.g., a third-party model vendor) – Ruth has no control over, and cannot indemnify for, a third party's own model or its outputs.
- **[Medium Risk]** A consequential-damages carve-back specifically for "loss of business opportunity" tied to system downtime, which in practice reintroduces effectively unlimited exposure through the back door.
- **[High Risk]** The liability cap and damages exclusions apply only to Ruth and not reciprocally to the client (a unilateral, non-mutual structure).

Negotiation Guidance

- Use a tiered cap structure as the standard negotiating framework: base cap = 12 months' fees; narrow carve-outs (IP infringement, confidentiality breach, gross negligence) at a higher defined multiple – but never present or accept an uncapped option.
- On AI-related liability specifically, insist that any indemnification obligation is limited to Ruth's own code and model IP, expressly excluding the behavior of third-party foundation models and the client's own use or deployment decisions.
- Offer increased Technology E&O or Cyber insurance coverage as the trade for a higher liability cap, rather than agreeing to remove the cap altogether.
- Always insist that the liability limitation and damages exclusions are mutual and reciprocal – a one-sided cap is a common and avoidable red flag in client-drafted paper.

9 SERVICE LEVELS AND PERFORMANCE STANDARDS

Preferred Position

- SLAs apply only to defined managed-services or support engagements, not to one-time build/development projects, with response times tiered by severity (e.g., Critical: 2 business hours; High: 8 business hours; Medium: 2 business days).
- Any uptime commitment applies only to infrastructure directly managed or hosted by Ruth, excludes third-party cloud or AI-provider outages, is measured monthly (e.g., 99.5% monthly uptime), and its sole remedy is a service credit – not damages – capped at 10% of the applicable monthly fees.
- For AI/ML systems specifically, "performance SLA" is defined narrowly as the availability and latency of the serving infrastructure; model output accuracy or quality is addressed through a periodic model performance review cadence (e.g., quarterly), not a penalty-bearing SLA.

Acceptable Alternatives

- Best-efforts (non-penalized) response times may apply for support engagements under USD 10,000 per month in recurring fees.
- SLA credits may escalate with repeated breaches (e.g., 5% of monthly fee per breach, capped at 20% total) rather than a flat single cap, for larger managed-services accounts where the client requires stronger incentives.
- A shared-responsibility uptime model may be used, jointly measured and excluding outages caused by the client (e.g., client-imposed API rate limits, client data pipeline failures).
- Model-quality "guardrail" metrics (e.g., an automated evaluation score must remain above a defined floor) may trigger a review-and-remediate obligation rather than a financial penalty, where the client requires some form of ongoing quality assurance.

High-Risk / Red Lines

- **[High Risk]** Any SLA or uptime commitment that includes third-party AI provider outages (e.g., an upstream model API provider) within Ruth's measured downtime – Ruth has no control over, and cannot be responsible for, an upstream provider's own SLA.
- **[High Risk]** SLA credits are not designated as the client's sole and exclusive remedy, allowing the client to pursue damages claims in addition to service credits for the same breach – this creates double exposure.
- **[Medium Risk]** No defined exclusions in the uptime calculation for scheduled maintenance windows, force majeure events, or client-caused outages.
- **[High Risk]** A binding SLA on AI model accuracy, precision, or recall in production, with financial penalties for missed thresholds, that does not account for data drift or changes in client-side data quality.
- **[Medium Risk]** Response-time SLA with no severity tiering, such that every issue – regardless of actual business impact – is treated as "critical."
- **[Medium Risk]** SLA service credits with no percentage cap relative to fees, exposing Ruth to disproportionate financial penalty relative to the underlying contract value.

Negotiation Guidance

- Clearly separate "infrastructure/availability SLA" (measurable, appropriately credit-bearing) from "model quality" (handled via a review cadence and improvement backlog, not a penalty-bearing SLA) – this distinction should be introduced early in any AI/ML SLA discussion.

- When a client insists on model performance commitments, propose "best efforts to maintain metric X on agreed test set Y, reviewed quarterly, with a defined remediation plan if the metric drifts below an agreed threshold" as a substitute for a hard penalty SLA.
- Always designate SLA credits as the sole and exclusive remedy for SLA breaches, explicitly excluding both termination-for-convenience and separate damages claims for the same underlying breach.
- Explicitly exclude force majeure events, third-party provider outages, and client-caused issues from the SLA measurement methodology in the SOW itself, not just in a general force majeure clause.

10 ACCEPTANCE TESTING AND DELIVERABLES

Preferred Position

- Deliverables are deemed accepted if the client does not provide written rejection, including specific and reproducible defects, within 10 business days of delivery (the "Acceptance Period").
- Acceptance criteria are defined in the SOW in advance – a functional specification, defined test cases, or, for AI/ML deliverables, an agreed evaluation dataset plus a minimum benchmark score – and no new or undocumented criteria may be introduced at the point of acceptance review.
- Where a deliverable is validly rejected, Ruth has a defined cure period (e.g., 10 business days) to remediate the documented defects, after which the deliverable is resubmitted and evaluated against the *same* originally agreed acceptance criteria.

Acceptable Alternatives

- An extended Acceptance Period (15–20 business days) may apply for complex, multi-module deliverables or for enterprise clients with a formal internal User Acceptance Testing (UAT) process.
- Partial acceptance may be permitted, allowing the client to accept non-defective modules while rejecting only the specific defective components, with payment released proportionally to accepted modules.
- Acceptance testing may be conducted jointly by Ruth and the client, rather than unilaterally by the client alone, to reduce disputes arising from undocumented or evolving criteria.
- For AI/ML deliverables, acceptance may be based on a statistical benchmark achieved on an agreed held-out test set, with a defined tolerance or confidence interval, rather than requiring an exact score match.

High-Risk / Red Lines

- **[High Risk]** No defined Acceptance Period at all, with deliverables "accepted upon client's sole discretion at any time" – this creates indefinite payment risk for Ruth with no contractual endpoint.
- **[High Risk]** Acceptance criteria are left undefined or stated only as "to the client's satisfaction," permitting subjective rejection with no objective recourse for Ruth.
- **[High Risk]** AI model acceptance is tied to a specific accuracy threshold measured against a test dataset that is not disclosed or agreed until *after* delivery – an unworkable moving target.
- **[Medium Risk]** Unlimited rejection-and-cure cycles with no cap on the number of resubmissions before a dispute-escalation or termination right is triggered.
- **[Medium Risk]** No partial acceptance mechanism, such that a single defective module blocks payment for the entire set of deliverables under a milestone.
- **[High Risk]** "Deemed rejection" (rather than deemed acceptance) applies if the client remains silent past the Acceptance Period – this reverses the default risk allocation onto Ruth indefinitely and should never be accepted.

Negotiation Guidance

- Insist on "deemed acceptance" language (client silence equals acceptance after the defined period) as a default – this is one of the most important protections in the entire playbook and should be held firm in negotiation.

- For AI/ML SOWs, lock the evaluation dataset and success metric into a signed SOW appendix *before* work begins; never leave benchmark data or methodology "to be determined" later in the engagement.
- Offer a reasonable, explicit cap on rejection-and-cure cycles (e.g., two cure cycles) before the parties move to the formal change-order or dispute-resolution process, to prevent open-ended rework at no additional charge.
- Use partial acceptance and milestone-based payment structures to keep cash flow moving even while a specific component is under active dispute or remediation.

11 TERM, TERMINATION AND EXIT ASSISTANCE

Preferred Position

- The initial term matches the applicable SOW duration; the MSA has a 12-month term, auto-renewing for successive 12-month periods unless either party gives 60 days' written notice of non-renewal.
- Either party may terminate for convenience with 30 days' written notice; upon such termination, the client pays for all work performed to date plus any non-cancellable committed costs (e.g., contracted staff already onboarded, third-party licenses purchased for the engagement) through the effective termination date.
- Termination for material breach requires written notice and a 15-business-day cure period before termination takes effect; immediate termination without a cure period is reserved for insolvency, unauthorized use of the other party's IP, or a material breach of confidentiality.

Acceptable Alternatives

- A client-only termination-for-convenience right with a longer notice period (45–60 days) may be agreed in exchange for no early-termination fee payable by Ruth, provided Ruth retains the right to redeploy affected staff freely.
- An early-termination fee (e.g., equivalent to one month's fees) may apply if the client terminates for convenience within the first 90 days of the engagement, to offset ramp-up and staffing costs already incurred.
- The cure period for payment defaults specifically may be shortened to 10 business days, while remaining at 15 business days for other categories of breach.
- A defined wind-down/transition period (15–30 days), billed at standard rates, may be built into every termination scenario instead of immediate cessation of services.

High-Risk / Red Lines

- **[High Risk]** Client may terminate for convenience with no notice period and no obligation to pay for work-in-progress or non-cancellable committed costs already incurred by Ruth.
- **[High Risk]** Termination automatically triggers an uncompensated transfer of all in-progress work product, model weights, and Ruth Background IP, including unfinished or unpaid milestones.
- **[Medium Risk]** No cure period is provided for any alleged breach, permitting immediate termination for even minor or disputed breaches.
- **[Medium Risk]** Non-renewal notice period is set at under 30 days, creating a risk of unplanned resource deallocation with insufficient lead time for Ruth's staffing plans.
- **[High Risk]** Exit-assistance and transition-out obligations are open-ended and unpaid (e.g., "Ruth shall provide all reasonable transition assistance at no charge for as long as required by the client").
- **[Medium Risk]** The survival clause fails to explicitly preserve payment obligations, confidentiality, IP ownership/licenses, and liability caps beyond termination, creating ambiguity about what actually survives.

Negotiation Guidance

- Always tie any termination-for-convenience right to payment for work performed plus non-cancellable committed costs – never accept a "no payment on termination" position, regardless of deal size.

- Cap free transition/exit assistance at a fixed period (e.g., 10 business days) and bill any exit assistance beyond that period at Ruth's standard T&M rates.
- Push for symmetrical (mutual) termination-for-convenience rights where possible – this is generally viewed as fairer by client legal teams and can speed up approval compared to a one-sided right.
- Explicitly enumerate the surviving clauses (payment obligations, IP terms, confidentiality, liability caps, indemnification) in a dedicated survival clause to eliminate ambiguity at termination.

12 KEY PERSONNEL, STAFFING AND SUBCONTRACTING

Preferred Position

- Named "Key Personnel" is limited to roles genuinely critical to the engagement (e.g., Engagement Lead, Lead ML Architect) – not the entire delivery team – and Ruth commits to 30 days' advance notice before any planned replacement of Key Personnel, along with reasonable knowledge transfer, except in cases of resignation, termination for cause, or illness beyond Ruth's control.
- Ruth retains sole discretion over general staffing, performance management, and internal team composition, subject to a defined client right to request replacement of a specific individual for documented performance or conduct issues, with a replacement provided within 10 business days.
- Subcontracting of work – including to Ruth affiliates or specialist contractors – is permitted without requiring client consent, provided Ruth remains fully responsible for the subcontractor's performance and for flow-down of confidentiality and IP obligations.

Acceptable Alternatives

- Client pre-approval may be required specifically for subcontractors who will have direct access to the client's production systems or sensitive data, without extending that requirement to general internal delivery subcontracting.
- The Key Personnel list may be limited to a maximum of one or two named individuals, with substitution permitted after a defined ramp/handover period requiring notice only, not client approval.
- A mutual non-solicitation clause covering Ruth's assigned personnel for 12 months post-engagement may be agreed, in exchange for Ruth committing not to reassign Key Personnel without advance notice.
- Backup or "shadow" resourcing (a secondary, cross-trained engineer) may be offered as a business-continuity measure in lieu of an absolute no-replacement guarantee for named individuals.

High-Risk / Red Lines

- **[High Risk]** Client holds a unilateral right to approve or reject every individual staffed on the project, including junior resources – unworkable at Ruth's 150-person scale and removes Ruth's core staffing autonomy.
- **[High Risk]** A Key Personnel clause where the departure of *any* listed individual (e.g., due to voluntary resignation) is itself defined as a termination trigger or contractual breach, regardless of cause.
- **[Medium Risk]** A client non-solicitation clause that is one-sided – restricting the client from hiring Ruth's staff while leaving Ruth free to hire the client's staff – should be made mutual.
- **[Medium Risk]** Subcontracting is prohibited entirely, including to Ruth's own global delivery centers or affiliates, conflicting directly with Ruth's distributed delivery model.
- **[High Risk]** Client requires that all AI/ML work be performed exclusively by personnel physically located in a specific country or region with no exceptions – may be a legitimate requirement for regulated industries, but must be flagged as a major staffing and cost constraint and priced accordingly before signing.
- **[Medium Risk]** No defined replacement timeline for a Key Personnel departure, leaving Ruth exposed to open-ended "immediate replacement" demands from the client.

Negotiation Guidance

- Offer a client escalation path (raising staffing concerns to the Engagement Lead or Delivery Director) as a workable alternative to individual approval rights over every team member.
- Keep the Key Personnel list short and tie any replacement commitment to notice-plus-transition, not a guarantee against natural staff attrition.
- If a client demands location-restricted staffing for data residency or regulatory reasons, quote it explicitly as a premium-priced option that reflects the smaller available resource pool, rather than absorbing the cost silently.
- Propose mutual non-solicitation as a standard counter to one-sided non-solicit demands – it protects the ongoing relationship while remaining fair to both parties.

13 OPEN SOURCE, THIRD-PARTY COMPONENTS AND AI TOOLING

Preferred Position

- Ruth may use open-source components and approved AI-based development tools (e.g., LLM-based coding assistants) in delivery, provided all open-source use complies with permissive or client-approved licenses (MIT, Apache 2.0, BSD), disclosed in a Bill of Materials / Software Bill of Materials (SBOM) delivered upon final acceptance.
- Use of copyleft-licensed components (GPL, AGPL, LGPL) requires the client's prior written approval before inclusion in any deliverable, given their potential to impose source-disclosure obligations on the client's own proprietary code.
- Ruth discloses which third-party or foundation AI models (e.g., commercial LLM APIs or open-weight models) are used within the solution architecture, including whether data is transmitted to an external API or processed via a self-hosted deployment, documented in the SOW or an accompanying technical architecture document.

Acceptable Alternatives

- A pre-approved standard technology stack list (including standard AI development/coding assistants) may be attached to the MSA and updated periodically, substituting for per-project disclosure of routine, well-known components.
- A client-specific "no GPL/AGPL" policy may be honored through an automated license-scanning report delivered at project completion, rather than a manual, individually-signed attestation.
- Where a client requires self-hosted or on-premises models only, with no external API calls, Ruth may agree subject to additional infrastructure fees and an extended delivery timeline.
- An SBOM may be delivered as a summary report (rather than a full version-level dependency tree) for lower-sensitivity engagements.

High-Risk / Red Lines

- **[High Risk]** Client contract is silent on open-source specifically but includes a broad "no third-party components without consent" clause that would, read literally, be violated by nearly all modern software development – must be flagged and negotiated, never signed as-is.
- **[High Risk]** Client prohibits all use of AI coding assistants or LLM-based tools in delivery, with no practical alternative offered, materially reducing delivery efficiency without a corresponding adjustment to price or timeline.
- **[Medium Risk]** Client requires a warranty that deliverables contain "no open-source software whatsoever" – unrealistic for modern software development and should be narrowed to a targeted restriction on unapproved copyleft licenses.
- **[High Risk]** Client requires Ruth to indemnify for IP or licensing issues arising from open-source components that were incorporated per the client's own specific instruction or mandate.
- **[Medium Risk]** No process is defined for handling newly discovered open-source vulnerabilities (e.g., CVEs) after delivery – should default to a reasonable-efforts patching commitment under a paid support contract, not an indefinite free obligation.
- **[High Risk]** Client demands visibility and control over which foundation AI model providers Ruth may use, while simultaneously holding Ruth liable for that provider's terms of service or model outputs, despite Ruth having no control over that third party.

Negotiation Guidance

- Proactively share Ruth's standard "Approved Tools & Open-Source Policy" document at project kickoff to preempt broad, unworkable blanket prohibitions before they get drafted into a contract.
- If a client insists on banning AI coding assistants, quantify the resulting timeline and cost impact and present it as a formal re-quote, rather than silently absorbing the resulting productivity loss.
- Always narrow blanket "no open source" demands down to a targeted copyleft-license restriction, which is typically the client's actual underlying concern in the vast majority of cases.
- Offer a standard SBOM or license-compliance report as a value-add deliverable to build client trust, without agreeing to unworkable absolute warranties.

14 INSURANCE, AUDIT RIGHTS AND COMPLIANCE

Preferred Position

- Ruth maintains Commercial General Liability insurance (minimum USD 1M per occurrence / USD 2M aggregate), Technology Errors & Omissions / Professional Liability (minimum USD 2M), Cyber Liability (minimum USD 2M), and statutory Workers' Compensation, with certificates of insurance provided to the client upon reasonable request.
- Client audit rights are limited to: (a) security and compliance questionnaires and, where available, sharing of SOC 2 / ISO 27001 reports; and (b) on-site or virtual audits no more than once every 12 months, with 30 days' advance written notice, conducted during business hours, and subject to confidentiality protections.
- Ruth complies with data protection laws applicable to the specific data processing activities under the engagement (e.g., GDPR, CCPA, India's DPDP Act, as relevant), as detailed in the applicable DPA, and will support reasonable client compliance requests (e.g., Data Protection Impact Assessments) beyond the agreed baseline scope at standard billing rates.

Acceptable Alternatives

- Increased insurance limits (e.g., Cyber Liability raised to USD 5M) may be agreed for engagements processing sensitive or regulated data, with the incremental cost priced into the engagement fee.
- Audit frequency may be increased to twice per 12 months for clients in regulated industries (financial services, healthcare), provided costs beyond a reasonable time allotment (e.g., 2 business days per audit) are billable to the client.
- A third-party independent auditor engaged at the client's cost may be used in place of the client's own internal audit team directly accessing Ruth's systems, where the client requires deeper technical verification.
- Compliance certifications not yet held (e.g., SOC 2 Type II, ISO 27001) may be pursued on a defined roadmap and timeline, transparently disclosed to the client at contract signature rather than promised as already in place.

High-Risk / Red Lines

- [High Risk] Unlimited audit rights (any time, any frequency, any scope – including source code and internal AI model weights or training pipelines belonging to unrelated clients) – this exposes Ruth's own confidential IP and other clients' confidential information.
- [High Risk] Insurance requirements set with no regard to policy limits realistically available or affordable at Ruth's scale (e.g., demanding USD 50M in cyber coverage for a mid-sized services firm).
- [Medium Risk] Audit clause requires Ruth to bear all costs of client-initiated audits, including travel and third-party consultant fees, with no reasonableness cap.
- [High Risk] Compliance obligation requires Ruth to warrant compliance with laws or regulatory regimes not disclosed or genuinely applicable to the actual jurisdiction or industry of the engagement (e.g., a healthcare-specific regulatory regime for a non-healthcare engagement).
- [Medium Risk] No confidentiality protection applies to information or systems disclosed during a client-initiated audit, creating a risk of inadvertently exposing other clients' data in a shared delivery environment.
- [High Risk] Client requires direct read access to Ruth's AI model training logs or pipelines that may contain other clients' proprietary data or model IP.

Negotiation Guidance

- Offer SOC 2 / ISO reports and completed security questionnaires as the primary compliance-verification mechanism, before agreeing to grant any on-site or system-level audit right.
- Cap audit frequency and duration explicitly in the contract, and require the client to bear the costs of any audit activity beyond the agreed standard scope.
- For regulated-industry engagements, proactively disclose Ruth's current certification status and roadmap rather than over-promising compliance capabilities Ruth does not yet hold.
- Insist that any audit involving code or systems takes place in a sandboxed or pre-reviewed environment, to prevent inadvertent exposure of other clients' confidential material during the audit.

15 GENERAL PROVISIONS (FORCE MAJEURE, GOVERNING LAW, ASSIGNMENT, MFN, NOTICES)

Preferred Position

- Force majeure covers events beyond a party's reasonable control – including natural disasters, war, pandemic, government action, internet or cloud infrastructure-wide outages, and platform-wide outages of third-party AI providers – excusing delay in performance (but not payment obligations for work already performed), with either party entitled to terminate if the event continues beyond 60 days.
- Governing law and venue default to Ruth's home jurisdiction, with disputes resolved first through good-faith negotiation, then through binding institutional arbitration, before either party may resort to litigation.
- No most-favored-nation (MFN) pricing commitments are granted by default; assignment of the agreement by either party requires the other party's prior written consent, except that Ruth may freely assign the agreement to an affiliate or in connection with a merger, acquisition, or sale of substantially all assets.

Acceptable Alternatives

- A neutral third-country governing law and arbitration seat (e.g., Singapore or the UK) may be agreed for large global enterprise clients who will not accept Ruth's home-jurisdiction governing law.
- Litigation (rather than arbitration) may be accepted for lower-value disputes (e.g., under USD 100,000), with arbitration reserved for disputes above that threshold.
- A narrow MFN commitment, limited to identical scope, volume, and term as compared against a specific, named benchmark client set, reviewed no more than once annually, may be agreed where commercially necessary to close a strategic deal.
- Assignment consent may be qualified as "not to be unreasonably withheld or delayed" rather than an absolute consent requirement, to ease routine corporate restructuring on either side.

High-Risk / Red Lines

- **[High Risk]** A broad MFN clause requiring Ruth to proactively disclose and match its lowest price given to *any* client globally – operationally unworkable and commercially damaging across a multi-client, 150-person services firm, and must not be agreed in broad form.
- **[High Risk]** Force majeure clause expressly excludes pandemics, government shutdowns, or third-party cloud/AI provider outages from covered events, leaving Ruth contractually liable for delays entirely outside its control.
- **[Medium Risk]** Governing law and venue set in a jurisdiction with no practical connection to either party and disproportionately high litigation cost (a forum-shopping risk).
- **[Medium Risk]** Assignment clause allows the client to freely assign the contract to any acquirer – including a direct competitor of Ruth – without Ruth's consent, while Ruth is required to obtain the client's consent to assign.
- **[High Risk]** Notices clause requires physical mail only, with no provision for email or electronic notice, creating operational risk of missed deadlines or notices in a fast-moving commercial relationship.
- **[Medium Risk]** The "entire agreement" clause omits, or contradicts, the DPA or security addendum, creating ambiguity about which data protection terms actually govern the relationship.

Negotiation Guidance

- If a client insists on an MFN clause, narrow its scope tightly (same service, same region, same volume tier) and limit any review to once per year, triggered only by written client request, not automatic audit.
- Always ensure third-party cloud and AI infrastructure outages are explicitly included within the force majeure definition, given Ruth's structural dependency on external AI and cloud providers for delivery.
- Propose institutional arbitration (under a recognized set of rules, e.g., SIAC, ICC, or LCIA) as a faster, more predictable middle-ground position when clients resist Ruth's home-jurisdiction litigation as the default.
- Make assignment rights mutual and add a "not to be unreasonably withheld" qualifier as the standard counter-offer, reducing friction with client legal teams while still protecting Ruth against undesirable assignees such as competitors.